

SENTIMENT ANALYSIS

Sentiment analysis is used to analyse what was said about a topic. Is the comment positive or negative? In this example we are using a data set of movie reviews. We need to classify each movie to positive/negative based on text given.

FEATURE EXTRACTION

The process of converting our text document to numerical form. I have used 3 types of vectorizers here. Vectorizers are evaluated based on the accuracies.

TF-IDF VECTORIZER

TF-IDF (Term frequency- Inverse Document Frequency) vectorizer considers the frequency of a word in a document and frequency between documents.

```
1 # training: tf-idf + Logistic regression
2 # you should explore different representations and algorithms.
3 from sklearn.feature_extraction.text import TfidfVectorizer
4 max_feature_num = 1000
5 train_vectorizer = TfidfVectorizer(max_features=max_feature_num)
6 train_vecs = train_vectorizer.fit_transform(train_text)
7 test_vecs = TfidfVectorizer(max_features=max_feature_num,vocabulary=train_vectorizer.vocabulary_).fit_transform(test_text)
8
9 # train model
10 from sklearn.linear_model import LogisticRegression
11 clf = LogisticRegression(max_iter = 5000).fit(train_vecs, train_labels)
12
13 # test model
14 test_pred = clf.predict(test_vecs)
15 from sklearn.metrics import precision_recall_fscore_support,accuracy_score
16 tfidf_acc = accuracy_score(test_labels, test_pred)
17 pre, rec, f1, _ = precision_recall_fscore_support(test_labels, test_pred, average='macro')
18 print('acc', tfidf_acc)
19 print('precision', pre)
20 print('rec', rec)
21 print('f1', f1)
```

```
acc 0.8616
precision 0.8616043455692743
rec 0.8615913854621674
f1 0.8615956596398864
```

COUNT VECTORIZER

Count vectorizer considers every word in the document and converts to features.

```
: 1 from sklearn.feature_extraction.text import CountVectorizer
2 max_feature_num = 1000
3 count_train_vectorizer = CountVectorizer(max_features=max_feature_num)
4 train_vecs = count_train_vectorizer.fit_transform(train_text)
5 test_vecs = CountVectorizer(max_features=max_feature_num,vocabulary=train_vectorizer.vocabulary_).fit_transform(test_text)
6
7 # train model
8 from sklearn.linear_model import LogisticRegression
9 clf = LogisticRegression(max_iter = 5000).fit(train_vecs, train_labels)
10
11 # test model
12 test_pred = clf.predict(test_vecs)
13 from sklearn.metrics import precision_recall_fscore_support,accuracy_score
14 count_acc = accuracy_score(test_labels, test_pred)
15 pre, rec, f1, _ = precision_recall_fscore_support(test_labels, test_pred, average='macro')
16 print('acc', count_acc)
17 print('precision', pre)
18 print('rec', rec)
19 print('f1', f1)
```

```
acc 0.8648
precision 0.8648402808583475
rec 0.8647786364581833
f1 0.8647895293011492
```

HASHING VECTORIZER

Converts a document a collection to matrix of word occurrences

```
: 1 from sklearn.feature_extraction.text import HashingVectorizer
2 max_feature_num = 1000
3 hash_train_vectorizer = HashingVectorizer(n_features=max_feature_num)
4 train_vecs = hash_train_vectorizer.fit_transform(train_text)
5 test_vecs = HashingVectorizer(n_features=max_feature_num).fit_transform(test_text)
6
7 # train model
8 from sklearn.linear_model import LogisticRegression
9 clf = LogisticRegression(max_iter = 5000).fit(train_vecs, train_labels)
10
11 # test model
12 test_pred = clf.predict(test_vecs)
13 from sklearn.metrics import precision_recall_fscore_support, accuracy_score
14 hash_acc = accuracy_score(test_labels, test_pred)
15 pre, rec, f1, _ = precision_recall_fscore_support(test_labels, test_pred, average='macro')
16 print('acc', hash_acc)
17 print('precision', pre)
18 print('rec', rec)
19 print('f1', f1)
```

```
acc 0.82
precision 0.820026833196351
rec 0.8199795196723147
f1 0.8199872983037684
```

COMPARING THE RESULTS OF VECTORIZERS

```
In [31]: 1 accuracy_df = pd.DataFrame()
2 accuracy_df['representation'] = ['tfidf_vec', 'count_vec', 'hash_vec']
3 accuracy_df['accuracy'] = [tfidf_acc, count_acc, hash_acc]
4 accuracy_df
5 |
6
```

Out[31]:

	representation	accuracy
0	tfidf_vec	0.8616
1	count_vec	0.8648
2	hash_vec	0.8200

Tfidf Vector and count vectorizer produced almost same result

We will be considering both count vectorizer and tfidf vectorizer for our next step.

Frequency distribution of words

TUNING HYPER PARAMETERS OF THE VECTORIZER

The frequency distribution of words in the data set is as follows


```
In [15]: 1 # frequency distribution of dataset
2 #making corpus containign all words of the dataset
3 import re
4 corpus = []
5 for sentence in all_text:
6     for word in sentence.split():
7         corpus.append(word.lower())
8
```

The above lines of code creates a list of all the unique words in the dataset.

```
In [16]: 1 #getting the frequency distribution of words in corpus
2 import nltk
3 freqdist = nltk.FreqDist(corpus)
4 print(freqdist.most_common(10))
5
```

[('the', 510245), ('a', 253107), ('and', 250840), ('of', 228709), ('to', 211262), ('is', 163617), ('in', 143713), ('i', 113318), ('this', 110472), ('it', 103506)]

```
1 freqdist.plot(100)
```



<matplotlib.axes._subplots.AxesSubplot at 0x2221d369b88>

Inorder to find out the best parameters for the vectorizer we use the GridSerachCV and Pipeline tool in sklearn.

Choosing Ngrams

```
5 from nltk.corpus import stopwords
6 #nltk.download('stopwords')
7 from nltk.stem import PorterStemmer
8 from sklearn.feature_extraction.text import TfidfVectorizer
9 from sklearn.linear_model import LogisticRegression
10
11 import re
12 def text_preprocess(text):
13     replace_punct = re.compile("[\,\.\(\)\[\]\,\;\:\!\?]") #punctuation removal
14     replace_tag = re.compile("(\\-|\\\/|<br\\s*\\\/>+)" #tags removal
15     text = [replace_punct.sub("", sent.lower()) for sent in text]
16     text = [replace_tag.sub(" ", line) for sent in text]
17     return text
18 def tokenize(all_text):
19     return all_text.split()
20 def norm_tokenize_stem(all_text):
21     porter = PorterStemmer()
22     for sent in all_text:
23         final_text = [porter.stem(w) for w in sent.split()]
24     return final_text
25 stop_words = stopwords.words('english')
26
27 vectorizer = TfidfVectorizer(strip_accents=None, lowercase=True)
28 clf = LogisticRegression(random_state=0,max_iter = 5000)
29 param_grid = [{'tfidf__ngram_range': [(1, 1),(1,2)],
30     'tfidf__stop_words': [stop_words,None],
31     'tfidf__tokenizer': [None,tokenize,norm_tokenize_stem],
32     'tfidf__preprocessor': [None,text_preprocess],
33     'tfidf__use_idf':[True,False],
34     'tfidf__max_features':[5000],
35     'tfidf__norm':['l2'],
36     'clf__C': [0.1,1.0,10.0]]]
37
38 pipe_tfidf_lr = Pipeline([('tfidf', vectorizer),
39     ('clf', clf)])
40 grid_tfidf_lr = GridSearchCV(pipe_tfidf_lr, param_grid, scoring='accuracy', cv=5, verbose=1, n_jobs=-1)
```

```
In [5]: 1 grid_tfidf_lr.fit(train_text,train_labels)
```

Fitting 5 folds for each of 144 candidates, totalling 720 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 34 tasks	elapsed: 2.7min
[Parallel(n_jobs=-1)]: Done 184 tasks	elapsed: 14.0min
[Parallel(n_jobs=-1)]: Done 434 tasks	elapsed: 28.8min
[Parallel(n_jobs=-1)]: Done 720 out of 720	elapsed: 43.4min finished

```
Out[5]: GridSearchCV(cv=5, error_score=nan,
    estimator=Pipeline(memory=None,
        steps=[('tfidf',
            TfidfVectorizer(analyzer='word',
                binary=False,
                decode_error='strict',
                dtype=<class 'numpy.float64'>,
                encoding='utf-8',
                input='content',
                lowercase=True,
                max_df=1.0,
                max_features=None,
                min_df=1,
                ngram_range=(1, 1),
                norm='l2',
                preprocessor=None,
                smooth_idf=True,
                stop_words=None,
                strip_acc...
                'yourselves', 'he', 'him',
                'his', 'himself', 'she',
                'she's', 'her', 'hers',
                'herself', 'it', 'it's', 'its',
                'itself', ...],
                None),
            'tfidf__tokenizer': [None,
                <function tokenize at 0x0000021AF33DC678>,
                <function norm_tokenize_stem at 0x0000021AF33E2168>],
```

```
In [6]: 1 print('Best parameters are: ' + str(grid_tfidf_lr.best_params_))
2 print('Best accuracy is : %.4f' % grid_tfidf_lr.best_score_)
```

Best parameters are: {'clf__C': 1.0, 'tfidf__max_features': 5000, 'tfidf__ngram_range': (1, 2), 'tfidf__norm': 'l2', 'tfidf__preprocessor': None, 'tfidf__stop_words': ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', 'she', "she's", 'her', 'hers', 'herself', 'it', 'it's', 'its', 'itself', 'they', 'them', 'their', 'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those', 'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', 'does', 'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', 'of', 'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', 'before', 'after', 'above', 'below', 'to', 'from', 'm', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under', 'again', 'further', 'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', 'few', 'more', 'most', 'other', 'some', 'such', 'no', 'non', 'not', 'only', 'own', 'same', 'so', 'than', 'too', 'very', 's', 't', 'can', 'will', 'just', 'don', 'don't', 'should', "should've", 'now', 'd', 'll', 'm', 'o', 're', 've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn', "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't", 'mustn', "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't", 'weren', "weren't", 'won', "won't", 'wouldn', "wouldn't"], 'tfidf__tokenizer': None, 'tfidf__use_idf': True}

Best accuracy is : 0.8845

- 1.The preprocess I used here will remove the hashtags and punctuations
- 2.Tokenizer used will split the word with space
3. Normalization used is Porter stemming.

The above results shows that:

1. Best value for C(logistic regression parameter) = 1.0
2. Best ngram_range = (1,2)
3. Preprocessing have no effect on getting best accuracy
4. Set of stopwords which improves the accuracy is ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', 'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their', 'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those', 'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', 'does', 'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', 'of', 'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', 'before', 'after', 'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under', 'again', 'further', 'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', 'few', 'more', 'most', 'other', 'some', 'such', 'no', 'nor', 'not', 'only', 'own', 'same', 'so', 'than', 'too', 'very', 's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', 'm', 'o', 're', 've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn', "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't", 'mustn', "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't", 'weren', "weren't", 'won', "won't", 'wouldn', "wouldn't"]
5. Tokenization have no effect on improving accuracy

In the following steps, we will be using the best parameters that we obtained from previous result for evaluating our vectorizers.

Apply the best parameters in the TF-IDF vectorizer and Count Vectorizer:

TF-IDF vectorizer

```
In [16]: 1 # training: tf-idf + Logistic regression
2 # you should explore different representations and algorithms.
3 from sklearn.feature_extraction.text import TfidfVectorizer
4 from nltk.tokenize import word_tokenize
5 from nltk.corpus import stopwords
6 max_feature_num = 5000
7 n_range = (1,2)
8 preprocess = None
9 stp_words = ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'you
10 nm = 'l2'
11 smoothIDF = True
12 token = None
13 useIDF = True
14 train_vectorizer = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words, ngram_range=n_range, max_f
15 train_vecs = train_vectorizer.fit_transform(train_text)
16 test_vecs = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words, ngram_range=n_range, max_features
17
18 # train model
19 from sklearn.linear_model import LogisticRegression
20 clf = LogisticRegression(C=1.0, max_iter = 5000).fit(train_vecs, train_labels)
21
22 # test model
23 test_pred = clf.predict(test_vecs)
24 from sklearn.metrics import precision_recall_fscore_support, accuracy_score
25 custom_tfidf_acc = accuracy_score(test_labels, test_pred)
26 pre, rec, f1, _ = precision_recall_fscore_support(test_labels, test_pred, average='macro')
27 print('acc', custom_tfidf_acc)
28 print('precision', pre)
29 print('rec', rec)
30 print('f1', f1)
31
32 <
33
34 acc 0.8868125
35 precision 0.8869615346072897
36 rec 0.8868125
37 f1 0.8868016006744337
```

Count vectorizer

```
In [17]: 1 # training: tf-idf + Logistic regression
2 # you should explore different representations and algorithms.
3 from sklearn.feature_extraction.text import CountVectorizer
4 from nltk.tokenize import word_tokenize
5 from nltk.corpus import stopwords
6 max_feature_num = 5000
7 n_range = (1,2)
8 preprocess = None
9 stp_words = ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'you
10 token = None
11 train_vectorizer = CountVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words, ngram_range=n_range, max_f
12 test_vecs = CountVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words, ngram_range=n_range, max_features
13
14 # train model
15 from sklearn.linear_model import LogisticRegression
16 clf = LogisticRegression(C=1.0, max_iter = 5000).fit(train_vecs, train_labels)
17
18 # test model
19 test_pred = clf.predict(test_vecs)
20 from sklearn.metrics import precision_recall_fscore_support, accuracy_score
21 custom_count_acc = accuracy_score(test_labels, test_pred)
22 pre, rec, f1, _ = precision_recall_fscore_support(test_labels, test_pred, average='macro')
23 print('acc', custom_count_acc)
24 print('precision', pre)
25 print('rec', rec)
26 print('f1', f1)
27
28 <
29
30 acc 0.6410625
31 precision 0.6503367982235213
32 rec 0.6410625
33 f1 0.6354400615324234
```

```
: 1 accuracy_df = pd.DataFrame()
2 accuracy_df['representation'] = ['tfidf_vec', 'count_vec']
3 accuracy_df['accuracy'] = [custom_tfidf_acc, custom_count_acc]
4 accuracy_df
5
```

```
:
   representation  accuracy
0      tfidf_vec  0.886813
1      count_vec  0.641062
```

The best params from grid search increased the accuracy of our sentiment classifier. But the accuracy of the count vectorizer has decreased.

Choosing maximum feature numbers

Trying different values for max_feature_num:

```
5 from nltk.corpus import stopwords
6 #nltk.download('stopwords')
7 from nltk.stem import PorterStemmer
8 from sklearn.feature_extraction.text import TfidfVectorizer
9 from sklearn.linear_model import LogisticRegression
10
11 import re
12 def text_preprocess(text):
13     replace_punct = re.compile("[\,\"()\\\[\]\.:\!\\'\"?]") #punctuation removal
14     replace_tag = re.compile("(\\-)|(\\/)|(\\<br\\s*>+)" #tags removal
15     text = [replace_punct.sub("", sent.lower()) for sent in text]
16     text = [replace_tag.sub(" ", line) for sent in text]
17     return text
18 def tokenize(all_text):
19     return all_text.split()
20 def norm_tokenize_stem(all_text):
21     porter = PorterStemmer()
22     for sent in all_text:
23         final_text = [porter.stem(w) for w in sent.split()]
24     return final_text
25 stop_words = stopwords.words('english')
26
27 vectorizer = TfidfVectorizer(strip_accents=None, lowercase=True)
28 clf = LogisticRegression(random_state=0,max_iter = 5000)
29 param_grid = [{'tfidf_ngram_range': [(1,2)],
30     'tfidf_stop_words': [stop_words,None],
31     'tfidf_tokenizer': [None],
32     'tfidf_preprocessor': [None],
33     'tfidf_use_idf':[True],
34     'tfidf_max_features':[5000,7000,9000,10000,12000,15000],
35     'tfidf_norm':['l2'],
36     'tfidf_smooth_idf':[True],
37     'clf_C': [1.0]}]
38
39 pipe_tfidf_lr = Pipeline([('tfidf', vectorizer),
40     ('clf', clf)])
41 grid_tfidf_lr = GridSearchCV(pipe_tfidf_lr, param_grid, scoring='accuracy', cv=5, verbose=1, n_jobs=-1)
```

```
5 from nltk.corpus import stopwords
6 #nltk.download('stopwords')
7 from nltk.stem import PorterStemmer
8 from sklearn.feature_extraction.text import TfidfVectorizer
9 from sklearn.linear_model import LogisticRegression
10
11 import re
12 def text_preprocess(text):
13     replace_punct = re.compile("[\,\"()\\\[\]\.:\!\\'\"?]") #punctuation removal
14     replace_tag = re.compile("(\\-)|(\\/)|(\\<br\\s*>+)" #tags removal
15     text = [replace_punct.sub("", sent.lower()) for sent in text]
16     text = [replace_tag.sub(" ", line) for sent in text]
17     return text
18 def tokenize(all_text):
19     return all_text.split()
20 def norm_tokenize_stem(all_text):
21     porter = PorterStemmer()
22     for sent in all_text:
23         final_text = [porter.stem(w) for w in sent.split()]
24     return final_text
25 stop_words = stopwords.words('english')
26
27 vectorizer = TfidfVectorizer(strip_accents=None, lowercase=True)
28 clf = LogisticRegression(random_state=0,max_iter = 5000)
29 param_grid = [{'tfidf_ngram_range': [(1,2)],
30     'tfidf_stop_words': [stop_words,None],
31     'tfidf_tokenizer': [None],
32     'tfidf_preprocessor': [None],
33     'tfidf_use_idf':[True],
34     'tfidf_max_features':[5000,7000,9000,10000,12000,15000],
35     'tfidf_norm':['l2'],
36     'tfidf_smooth_idf':[True],
37     'clf_C': [1.0]}]
38
39 pipe_tfidf_lr = Pipeline([('tfidf', vectorizer),
40     ('clf', clf)])
41 grid_tfidf_lr = GridSearchCV(pipe_tfidf_lr, param_grid, scoring='accuracy', cv=5, verbose=1, n_jobs=-1)
```

scoring: accuracy, verbose: 2,

```
In [96]: 1 print('Best parameter set: ' + str(grid_tfidf_lr.best_params_))
2 print('Best accuracy: %.3f' % grid_tfidf_lr.best_score_)
```

```
Best parameter set: {'clf_C': 1.0, 'tfidf_max_features': 15000, 'tfidf_ngram_range': (1, 2), 'tfidf_norm': 'l2', 'tfidf_preprocessor': None, 'tfidf_smooth_idf': True, 'tfidf_stop_words': None, 'tfidf_tokenizer': None, 'tfidf_use_idf': True}
Best accuracy: 0.889
```

The number of features which gives the maximum accuracy is 15000. Train the data set in combination with `max_features = 15000` and the other params we choose earlier gives:

```
In [19]: 1 # training: tf-idf + Logistic regression
2 # you should explore different representations and algorithms.
3 from sklearn.feature_extraction.text import TfidfVectorizer
4 from nltk.tokenize import word_tokenize
5 from nltk.corpus import stopwords
6 max_feature_num = 15000
7 n_range = (1,2)
8 preprocess = None
9 stp_words = ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'you
10 nm = 'l2'
11 smoothIDF = True
12 token = None
13 useIDF = True
14 train_vectorizer = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words, ngram_range=n_range, max_f
15 train_vecs = train_vectorizer.fit_transform(train_text)
16 test_vecs = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words, ngram_range=n_range, max_features
17 |
18 # train model
19 from sklearn.linear_model import LogisticRegression
20 clf = LogisticRegression(C=1.0, max_iter = 5000).fit(train_vecs, train_labels)
21
22 # test model
23 test_pred = clf.predict(test_vecs)
24 from sklearn.metrics import precision_recall_fscore_support, accuracy_score
25 custom_tfidf_acc = accuracy_score(test_labels, test_pred)
26 pre, rec, f1, _ = precision_recall_fscore_support(test_labels, test_pred, average='macro')
27 print('acc', custom_tfidf_acc)
28 print('precision', pre)
29 print('rec', rec)
30 print('f1', f1)
31
32 <
33 >
```

acc 0.8901875
precision 0.8903340277346301
rec 0.8901874999999999
f1 0.8901771933869965

The accuracy now has improved to 0.890175

Tuning max_df

```
25 stop_words = stopwords.words('english')
26 stp_words = ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'you
27
28 vectorizer = TfidfVectorizer(strip_accents=None, lowercase=True)
29 clf = LogisticRegression(random_state=0, max_iter = 5000)
30 param_grid = [{'tfidf__ngram_range': [(1,2)],
31 'tfidf__max_df': [0.25, 0.5, 0.75, 1.0],
32 'tfidf__stop_words': [stp_words, None],
33 'tfidf__tokenizer': [None],
34 'tfidf__preprocessor': [None],
35 'tfidf__use_idf': [True],
36 'tfidf__max_features': [15000],
37 'tfidf__norm': ['l2'],
38 'tfidf__smooth_idf': [True],
39 'clf__C': [1.0]]}
40
41 pipe_tfidf_lr = Pipeline([('tfidf', vectorizer),
42 ('clf', clf)])
43 grid_tfidf_lr = GridSearchCV(pipe_tfidf_lr, param_grid, scoring='accuracy', cv=5, verbose=1, n_jobs=-1)
44
45 <
46 >
```


Accuracy has now decreased to .8895

Tuning logistic regression params

```
22     for sent in all_text:
23         final_text = [porter.stem(w) for w in sent.split()]
24     return final_text
25 stop_words = stopwords.words('english')
26 stp_words = ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'you
27
28 vectorizer = TfidfVectorizer(strip_accents=None, lowercase=True)
29 clf = LogisticRegression(random_state=0, max_iter=5000)
30 param_grid = [{'tfidf__ngram_range': [(1, 2)],
31         'tfidf__max_df': [0.25],
32         'tfidf__stop_words': [stp_words, None],
33         'tfidf__tokenizer': [None],
34         'tfidf__preprocessor': [None],
35         'tfidf__use_idf': [True],
36         'tfidf__max_features': [15000],
37         'tfidf__norm': ['l2'],
38         'tfidf__smooth_idf': [True],
39         'clf__C': [0.1, 0.5, 1.0, 10.0, 100.0]}]
40
41 pipe_tfidf_lr = Pipeline(['tfidf', vectorizer],
42        ('clf', clf))
43 grid_tfidf_lr = GridSearchCV(pipe_tfidf_lr, param_grid, scoring='accuracy', cv=5, verbose=1, n_jobs=-1)
```

```
In [43]: 1 grid_tfidf_lr.fit(train_text, train_labels)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks | elapsed: 5.8min
[Parallel(n_jobs=-1)]: Done 50 out of 50 | elapsed: 8.1min finished
```

```
Out[43]: GridSearchCV(cv=5, error_score=nan,
        estimator=Pipeline(memory=None,
        steps=[('tfidf',
        TfidfVectorizer(analyzer='word',
        binary=False,
        decode_error='strict',
        dtype=<class 'numpy.float64'>,
        encoding='utf-8',
        input='content',
        lowercase=True,
        max_df=1.0,
        max_features=None,
        min_df=1,
        ngram_range=(1, 1),
        norm='l2',
        preprocessor=None,
        smooth_idf=True,
        stop_words=None,
        strip_acc...
        'we', 'our', 'ours',
        'ourselves', 'you', "you're",
        "you've", "you'll", "you'd",
        'your', 'yours', 'yourself',
        'yourselves', 'he', 'him',
        'his', 'himself', 'she',
        "she's", 'her', 'hers',
        'herself', 'it', "it's", 'its',
        '...')])
```

```
: 1 print('Best parameter set: ' + str(grid_tfidf_lr.best_params_))
2 print('Best accuracy: %.3f' % grid_tfidf_lr.best_score_)
```

```
Best parameter set: {'clf__C': 1.0, 'tfidf__max_df': 0.25, 'tfidf__max_features': 15000, 'tfidf__ngram_range': (1, 2), 'tfidf__
norm': 'l2', 'tfidf__preprocessor': None, 'tfidf__smooth_idf': True, 'tfidf__stop_words': None, 'tfidf__tokenizer': None, 'tfidf
f__use_idf': True}
Best accuracy: 0.898
```

The best value of C is 1.00

Tuning the penalty:

```
24 from sklearn.feature_extraction.text import TfidfVectorizer
25 stop_words = stopwords.words('english')
26 stp_words = ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll", "you'd", 'you
27
28 vectorizer = TfidfVectorizer(strip_accents=None, lowercase=True)
29 clf = LogisticRegression(random_state=0, max_iter=5000)
30 param_grid = [{'tfidf__ngram_range': [(1,2)],
31                  'tfidf__max_df': [0.25],
32                  'tfidf__stop_words': [stp_words, None],
33                  'tfidf__tokenizer': [None],
34                  'tfidf__preprocessor': [None],
35                  'tfidf__use_idf': [True],
36                  'tfidf__max_features': [15000],
37                  'tfidf__norm': ['l2'],
38                  'tfidf__smooth_idf': [True],
39                  'clf__penalty': ['l1', 'l2'],
40                  'clf__C': [1.0],
41                  'clf__solver': ['liblinear']}]
42
43 pipe_tfidf_lr = Pipeline([('tfidf', vectorizer),
44                           ('clf', clf)])
45 grid_tfidf_lr = GridSearchCV(pipe_tfidf_lr, param_grid, scoring='accuracy', cv=5, verbose=1, n_jobs=-1)
```

In [48]: 1 grid_tfidf_lr.fit(train_text, train_labels)

Fitting 5 folds for each of 4 candidates, totalling 20 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 20 out of 20 | elapsed: 3.2min finished

Out[48]: GridSearchCV(cv=5, error_score=nan,
 estimator=Pipeline(memory=None,
 steps=[('tfidf',
 TfidfVectorizer(analyzer='word',
 binary=False,
 decode_error='strict',
 dtype=<class 'numpy.float64'>,
 encoding='utf-8',
 input='content',
 lowercase=True,
 max_df=1.0,
 max_features=None,
 min_df=1,
 ngram_range=(1, 1).
 stop_words=None,
 tokenizer=None,
 transformer_decoder=None,
 transformer_encoder=None,
 use_idf=True,
 verbose=0,
 wrap_feature=0)],
 scoring='accuracy', verbose=1)

In [49]: 1 print('Best parameter set: ' + str(grid_tfidf_lr.best_params_))
 2 print('Best accuracy: %.3f' % grid_tfidf_lr.best_score_)

Best parameter set: {'clf__C': 1.0, 'clf__penalty': 'l2', 'clf__solver': 'liblinear', 'tfidf__max_df': 0.25, 'tfidf__max_features': 15000, 'tfidf__ngram_range': (1, 2), 'tfidf__norm': 'l2', 'tfidf__preprocessor': None, 'tfidf__smooth_idf': True, 'tfidf__stop_words': None, 'tfidf__tokenizer': None, 'tfidf__use_idf': True}
Best accuracy: 0.898

```

In [50]: 1 # training: tf-idf + logistic regression
2 # you should explore different representations and algorithms.
3 from sklearn.feature_extraction.text import TfidfVectorizer
4 from nltk.tokenize import word_tokenize
5 from nltk.corpus import stopwords
6 max_feature_num = 15000
7 n_range = (1,2)
8 preprocess = None
9 stop_words = None
10
11 smoothIDF = True
12 token = None
13 useIDF = True
14 maxdf = 0.25
15 train_vectorizer = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stop_words, ngram_range=n_range, max_df=maxdf)
16 train_vecs = train_vectorizer.fit_transform(train_text)
17 test_vecs = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stop_words, ngram_range=n_range, max_df=maxdf).fit_transform(test_text)
18
19 # train model
20 from sklearn.linear_model import LogisticRegression
21 clf = LogisticRegression(penalty='l2', C=1.0, solver='liblinear', max_iter=5000).fit(train_vecs, train_labels)
22 stop_words = None
23 nm = 'l2'
24 # test model
25 test_pred = clf.predict(test_vecs)
26 from sklearn.metrics import precision_recall_fscore_support, accuracy_score
27 custom_tfidf_acc = accuracy_score(test_labels, test_pred)
28 pre, rec, f1, _ = precision_recall_fscore_support(test_labels, test_pred, average='macro')
29 print('acc', custom_tfidf_acc)
30 print('precision', pre)
31 print('rec', rec)
32 print('f1', f1)

```

acc 0.8994
 precision 0.8994578972578153
 rec 0.8993767900286405
 f1 0.8993911101984972

The solver used here is liblinear and penalty is 'l2'(Ridge regression)

Comparing all best params we obtain from above:

50:50 Train/test split with test_size 0.3(balanced classes)

```

7
8 from sklearn.model_selection import train_test_split
9
10 train_text, test_text, train_labels, test_labels = train_test_split(all_text, all_labels, test_size=0.3, random_state=0, stratify=all_labels)
11

```

Here the parameter stratify make the number of pos and neg classes equal.

```

1 accuracy_list = [acc, best_ngram_acc, best_num_features_acc, best_maxdf_acc, best_penalty_acc]
2 best_params = ['acc', 'best_ngram_acc', 'best_num_features_acc', 'best_maxdf_acc', 'best_penalty_acc']
3 accuracy_df = pd.DataFrame()
4 accuracy_df['best_params'] = best_params
5 accuracy_df['accuracy'] = accuracy_list
6 accuracy_df.sort_values('accuracy', ascending=False)

```

	best_params	accuracy
3	best_maxdf_acc	0.900750
4	best_penalty_acc	0.900750
2	best_num_features_acc	0.894000
1	best_ngram_acc	0.888500
0	acc	0.866583

50:50 train/test split with test_size 0.4(balanced classes)

```
from sklearn.model_selection import train_test_split

train_text, test_text, train_labels, test_labels = train_test_split(all_text, all_labels, test_size=0.4, random_state=0, str
```

```
: 1 accuracy_list = [acc,best_ngram_acc, best_num_features_acc, best_maxdf_acc,best_penalty_acc]
2 best_params = ['acc','best_ngram_acc', 'best_num_features_acc', 'best_maxdf_acc','best_penalty_acc']
3 accuracy_df = pd.DataFrame()
4 accuracy_df['best_params'] = ['acc','best_ngram_acc', 'best_num_features_acc', 'best_maxdf_acc','best_penalty_acc']
5 accuracy_df['accuracy'] = accuracy_list
6 accuracy_df.sort_values('accuracy', ascending=False)
```


```
:
      best_params  accuracy
3  best_maxdf_acc  0.898125
4  best_penalty_acc  0.898125
2  best_num_features_acc  0.890188
1  best_ngram_acc  0.886813
0      acc  0.864125
```

Train/test split already given:

```
3 train_text = all_text[:35000]
4 train_labels = all_labels[:35000]
5 test_text = all_text[35000:]
6 test_labels = all_labels[35000:]
7
```

Comparing different train/test split ratios

```
In [33]: 1 accuracy_list_03 = [0.8665833333333334, 0.8885, 0.894, 0.90075, 0.90075]
2 accuracy_list_04 = [0.864125, 0.8868125, 0.8901875, 0.898125, 0.898125]
3 accuracy_initial = [0.8616, 0.8876, 0.8932, 0.8994, 0.8994]
4 best_params = ['acc','best_ngram_acc', 'best_num_features_acc', 'best_maxdf_acc','best_penalty_acc']
5 plt.figure()
6 plt.plot(best_params,accuracy_list_03,label = 'test size = 0.3',zorder = 1)
7 plt.plot(best_params,accuracy_list_04,label = 'test size = 0.4',zorder = 2)
8 plt.plot(best_params,accuracy_initial,label = 'train/test initial',zorder = 10)
9 plt.xticks(rotation = 45)
10 plt.show()
```



test_size = 0.3 shows maximum accuracy

SGD based sentiment analysis: (Feature based method)

SGD classifiers are classifiers with with a stochastic gradient descent (SGD) learning for gradient loss.

```

4 from nltk.tokenize import word_tokenize
5 from nltk.corpus import stopwords
6 max_feature_num = 15000
7 n_range = (1,2)
8 preprocess = None
9 stp_words = None
10
11 smoothIDF = True
12 token = None
13 useIDF = True
14 maxdf = 0.25
15 train_vectorizer = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words,ngram_range=n_range,max_df
16 train_vecs = train_vectorizer.fit_transform(train_text)
17 test_vecs = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words,ngram_range=n_range,max_df = maxd
18
19 # train model
20 from sklearn.linear_model import SGDClassifier
21 clf = SGDClassifier(loss='hinge', penalty='l2', alpha=1e-3, random_state=0, max_iter=100,learning_rate='optimal',tol=None).f
22 stp_words = None
23 nm = 'l2'
24 # test model
25 test_pred = clf.predict(test_vecs)
26 from sklearn.metrics import precision_recall_fscore_support,accuracy_score
27 best_penalty_acc = accuracy_score(test_labels, test_pred)
28 pre, rec, f1, _ = precision_recall_fscore_support(test_labels, test_pred, average='macro')
29 print('acc', best_penalty_acc)
30 print('precision', pre)
31 print('rec', rec)
32 print('f1', f1)

```

acc 0.8679166666666667
precision 0.8713750035434138
rec 0.8679166666666667
f1 0.8676084508658104

Decision Tree Based Sentiment analysis:

```

4 from nltk.tokenize import word_tokenize
5 from nltk.corpus import stopwords
6 max_feature_num = 15000
7 n_range = (1,2)
8 preprocess = None
9 stp_words = None
10
11 smoothIDF = True
12 token = None
13 useIDF = True
14 maxdf = 0.25
15 train_vectorizer = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words,ngram_range=n_range,max_df
16 train_vecs = train_vectorizer.fit_transform(train_text)
17 test_vecs = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words,ngram_range=n_range,max_df = maxd
18
19 # train model
20 from sklearn.tree import DecisionTreeClassifier
21 clf = DecisionTreeClassifier(criterion='entropy', random_state=0).fit(train_vecs, train_labels)
22 stp_words = None
23 nm = 'l2'
24 # test model
25 test_pred = clf.predict(test_vecs)
26 from sklearn.metrics import precision_recall_fscore_support,accuracy_score
27 dec_tree_accuracy = accuracy_score(test_labels, test_pred)
28 pre, rec, f1, _ = precision_recall_fscore_support(test_labels, test_pred, average='macro')
29 print('acc', dec_tree_accuracy)
30 print('precision', pre)
31 print('rec', rec)
32 print('f1', f1)

```

acc 0.7248333333333333
precision 0.7248335581668914
rec 0.7248333333333334
f1 0.7248332645416495

Naïve Bayes based sentiment analysis:

```
4 from nltk.tokenize import word_tokenize
5 from nltk.corpus import stopwords
6 max_feature_num = 15000
7 n_range = (1,2)
8 preprocess = None
9 stp_words = None
10
11 smoothIDF = True
12 token = None
13 useIDF = True
14 maxdf = 0.25
15 train_vectorizer = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words,ngram_range=n_range,max_df=
16 train_vecs = train_vectorizer.fit_transform(train_text)
17 test_vecs = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words,ngram_range=n_range,max_df = maxd
18
19 # train model
20 from sklearn.naive_bayes import MultinomialNB
21 clf = MultinomialNB().fit(train_vecs, train_labels)
22 stp_words = None
23 nm = 'l2'
24 # test model
25 test_pred = clf.predict(test_vecs)
26 from sklearn.metrics import precision_recall_fscore_support,accuracy_score
27 MNB_accuracy = accuracy_score(test_labels, test_pred)
28 pre, rec, f1, _ = precision_recall_fscore_support(test_labels, test_pred, average='macro')
29 print('acc', MNB_accuracy)
30 print('precision', pre)
31 print('rec', rec)
32 print('f1', f1)
```

< >

acc 0.8681666666666666
precision 0.8682532470967619
rec 0.8681666666666668
f1 0.8681589173408082

SVC based sentiment analysis:

```
3 from sklearn.feature_extraction.text import TfidfVectorizer
4 from nltk.tokenize import word_tokenize
5 from nltk.corpus import stopwords
6 max_feature_num = 15000
7 n_range = (1,2)
8 preprocess = None
9 stp_words = None
10
11 smoothIDF = True
12 token = None
13 useIDF = True
14 maxdf = 0.25
15 train_vectorizer = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words,ngram_range=n_range,max_df=
16 train_vecs = train_vectorizer.fit_transform(train_text)
17 test_vecs = TfidfVectorizer(preprocessor=preprocess, tokenizer=token, stop_words=stp_words,ngram_range=n_range,max_df = maxd
18
19 # train model
20 from sklearn.svm import SVC
21 clf = SVC(kernel='linear').fit(train_vecs, train_labels)
22 stp_words = None
23 nm = 'l2'
24 # test model
25 test_pred = clf.predict(test_vecs)
26 from sklearn.metrics import precision_recall_fscore_support,accuracy_score
27 SVC_accuracy = accuracy_score(test_labels, test_pred)
28 pre, rec, f1, _ = precision_recall_fscore_support(test_labels, test_pred, average='macro')
29 print('acc', SVC_accuracy)
30 print('precision', pre)
31 print('rec', rec)
32 print('f1', f1)
```

< >

acc 0.8975
precision 0.8975481033205017
rec 0.8975
f1 0.8974968992812034

Comparing the accuracy of different algorithms

```
1 accuracy_df = pd.DataFrame()
2 accuracy_df['model'] = ['best_lr_accuracy', 'sgd_accuracy', 'dec_tree_accuracy', 'SVC_accuracy', 'MNB_accuracy']
3 accuracy_df['accuracy'] = [best_lr,sgd_accuracy,dec_tree_accuracy,SVC_accuracy,MNB_accuracy]
4 accuracy_df.sort_values('accuracy',ascending = False)
5 accuracy_df
```

	model	accuracy
0	best_lr_accuracy	0.900750
1	sgd_accuracy	0.867917
2	dec_tree_accuracy	0.724833
3	SVC_accuracy	0.897500
4	MNB_accuracy	0.868167

Logistic Regression performs better than all other algorithms

So the final model I chose is TF-IDF vectorization with Logistic regression , 50:50 train test split(balanced classes),max_features = 15000,max_df = 0.25,and C= 1.00, penalty = 'l2' and solver = 'liblinear'

#MarkingId:25309